

```
from torchvision import datasets
from torchvision.transforms import ToTensor
from torch.utils.data import DataLoader
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torch

# Load MNIST dataset for training and testing
train_data = datasets.MNIST(
    root='data',          # Directory to store the dataset
    train=True,           # Load the training set
    transform=ToTensor(), # Apply ToTensor image transformation
    download=True         # Download the dataset if not already present
)

test_data = datasets.MNIST(
    root='data',
    train=False,          # Load the test set
    transform=ToTensor(),
    download=True
)

# Check the size of the datasets (number of samples, width, height)
print(train_data.data.size()) # Output: torch.Size([60000, 28, 28])
print(test_data.data.size())  # Output: torch.Size([10000, 28, 28])

# Check the labels for the training data
print(train_data.targets)
```

🔄 Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>
Failed to download (trying next):
HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz>
Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz>
100%|██████████| 9912422/9912422 [00:00<00:00, 13167454.61it/s]
Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>
Failed to download (trying next):
HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz>
Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz>
100%|██████████| 28881/28881 [00:00<00:00, 394998.25it/s]
Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>
Failed to download (trying next):

HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz>
Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz> t
100%|██████████| 1648877/1648877 [00:00<00:00, 3437466.99it/s]
Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>
Failed to download (trying next):
HTTP Error 403: Forbidden

Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz>
Downloading <https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz> t
100%|██████████| 4542/4542 [00:00<00:00, 9937678.02it/s]Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

```
torch.Size([60000, 28, 28])  
torch.Size([10000, 28, 28])  
tensor([5, 0, 4, ..., 5, 6, 8])
```

```
# Create data loaders for training and test sets
```

```
loader = {  
    'train': DataLoader(train_data,  
                        batch_size=125, # Number of samples per batch  
                        shuffle=True,   # Shuffle the data at every epoch  
                        num_workers=1), # Use 1 worker for data loading  
    'test': DataLoader(test_data,  
                      batch_size=125, # Number of samples per batch  
                      shuffle=False,  # No need to shuffle test data  
                      num_workers=1)  # Use 1 worker for data loading  
}
```

```
# Define the Convolutional NN model class
```

```
class CNN(nn.Module):  
    def __init__(self):  
        super(CNN, self).__init__()  
        # First convolutional layer: 1 input channel (grayscale), 10 output channels, kernel size 5  
        self.conv1 = nn.Conv2d(1, 10, kernel_size=5)  
  
        # Second convolutional layer: 10 input channels, 20 output channels, kernel size 5  
        self.conv2 = nn.Conv2d(10, 20, kernel_size=5)  
  
        # Dropout layer after the second convolution  
        self.conv2_dropout = nn.Dropout2d()  
  
        # Fully connected layer: 320 input features, 50 output features  
        self.fc1 = nn.Linear(320, 50)  
  
        # Fully connected output layer: 50 input features, 10 output features (10 digits)  
        self.fc2 = nn.Linear(50, 10)
```

```
def forward(self, x):
    # Apply first convolution, followed by ReLU and max pooling
    x = F.relu(F.max_pool2d(self.conv1(x), 2))

    # Apply second convolution, followed by dropout, ReLU, and max pooling
    x = F.relu(F.max_pool2d(self.conv2_dropout(self.conv2(x)), 2))

    # Flatten the tensor into a 1D vector (for input to fully connected layer)
    x = x.view(-1, 320)

    # Apply first fully connected layer followed by ReLU
    x = F.relu(self.fc1(x))

    # Apply dropout (during training only)
    x = F.dropout(x, training=self.training)

    # Apply second fully connected layer (output layer)
    x = self.fc2(x)

    # Use log_softmax as the final activation since we're using CrossEntropyLoss
    return F.log_softmax(x, dim=1)

# Set the device to use (GPU if available, otherwise CPU)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

# Initialize the model and move it to the appropriate device (CPU or GPU)
model = CNN().to(device)

# Define the optimizer (Adam) and the learning rate
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Define the loss function (CrossEntropyLoss, which combines softmax and negative log lik
loss_fn = nn.CrossEntropyLoss()

# Define the training loop
def train(epoch):
    model.train() # Set the model to training mode
    for batch_idx, (data, target) in enumerate(loader['train']):
        data, target = data.to(device), target.to(device) # Move data to the correct dev
        optimizer.zero_grad() # Zero out gradients from previous step
        output = model(data) # Forward pass
        loss = loss_fn(output, target) # Compute the loss
        loss.backward() # Backpropagate the gradients
        optimizer.step() # Update the model's weights
        if batch_idx % 20 == 0: # Print progress every 20 batches
            print(f'Train epoch: {epoch} [{batch_idx * len(data)} / {len(loader["train"]}
```

```

# Define the testing loop
def test():
    model.eval() # Set the model to evaluation mode
    test_loss = 0
    correct = 0

    with torch.no_grad(): # Disable gradient calculation for evaluation
        for data, target in loader['test']:
            data, target = data.to(device), target.to(device) # Move data to device
            output = model(data) # Forward pass
            test_loss += loss_fn(output, target).item() # Accumulate test loss
            pred = output.argmax(dim=1, keepdim=True) # Get the index of the max log-probability
            correct += pred.eq(target.view_as(pred)).sum().item() # Count correct predictions

    avg_loss = test_loss / len(loader['test']) # Compute average loss per batch
    accuracy = 100 * correct / len(loader['test'].dataset) # Compute accuracy percentage
    print(f'\nTest set: Average loss: {avg_loss:.4f}, Accuracy: {correct} / {len(loader["test"].dataset)}')

# Run the training and testing loops for 10 epochs
for epoch in range(1, 11):
    train(epoch) # Train the model for one epoch
    test() # Test the model after each epoch

```

```

Train epoch: 1 [0 / 60000]
Train epoch: 1 [2500 / 60000]
Train epoch: 1 [5000 / 60000]
Train epoch: 1 [7500 / 60000]
Train epoch: 1 [10000 / 60000]
Train epoch: 1 [12500 / 60000]
Train epoch: 1 [15000 / 60000]
Train epoch: 1 [17500 / 60000]
Train epoch: 1 [20000 / 60000]
Train epoch: 1 [22500 / 60000]
Train epoch: 1 [25000 / 60000]
Train epoch: 1 [27500 / 60000]
Train epoch: 1 [30000 / 60000]
Train epoch: 1 [32500 / 60000]
Train epoch: 1 [35000 / 60000]
Train epoch: 1 [37500 / 60000]
Train epoch: 1 [40000 / 60000]
Train epoch: 1 [42500 / 60000]
Train epoch: 1 [45000 / 60000]
Train epoch: 1 [47500 / 60000]
Train epoch: 1 [50000 / 60000]
Train epoch: 1 [52500 / 60000]
Train epoch: 1 [55000 / 60000]
Train epoch: 1 [57500 / 60000]

```

```
Train epoch: 1 [57500 / 60000]
```

```
Test set: Average loss: 0.1631, Accuracy: 9486 / 10000, 95%
```

```
Train epoch: 2 [0 / 60000]
```

```
Train epoch: 2 [2500 / 60000]
```

```
Train epoch: 2 [5000 / 60000]
```

```
Train epoch: 2 [7500 / 60000]
```

```
Train epoch: 2 [10000 / 60000]
```

```
Train epoch: 2 [12500 / 60000]
```

```
Train epoch: 2 [15000 / 60000]
```

```
Train epoch: 2 [17500 / 60000]
```

```
Train epoch: 2 [20000 / 60000]
```

```
Train epoch: 2 [22500 / 60000]
```

```
Train epoch: 2 [25000 / 60000]
```

```
Train epoch: 2 [27500 / 60000]
```

```
Train epoch: 2 [30000 / 60000]
```

```
Train epoch: 2 [32500 / 60000]
```

```
Train epoch: 2 [35000 / 60000]
```

```
Train epoch: 2 [37500 / 60000]
```

```
Train epoch: 2 [40000 / 60000]
```

```
Train epoch: 2 [42500 / 60000]
```

```
Train epoch: 2 [45000 / 60000]
```

```
Train epoch: 2 [47500 / 60000]
```

```
Train epoch: 2 [50000 / 60000]
```

```
Train epoch: 2 [52500 / 60000]
```

```
Train epoch: 2 [55000 / 60000]
```

```
Train epoch: 2 [57500 / 60000]
```

```
Test set: Average loss: 0.1063, Accuracy: 9656 / 10000, 97%
```

```
Train epoch: 3 [0 / 60000]
```

```
Train epoch: 3 [2500 / 60000]
```

```
Train epoch: 3 [5000 / 60000]
```

```
Train epoch: 3 [7500 / 60000]
```

